

Evidence for the Twelve Decisions

DAMM v1.7 · 24 August 2026 · Randeep Sudan

1. What Was Run, and How

Three exercises were carried out between 23 and 24 August 2026, none of which existed when the review package was assembled.

The vendor audition. Standing decision 4 fixed a method for choosing a research model and the method had never been run. Thirteen cells: ten whose answers are recorded in the Egypt and Nigeria assessments, and three naming statistics those assessments record as data gaps after a documented tier-ordered search. Five models from three vendors were given an identical set of fetched pages for each cell and scored on fabrication rate, tier compliance, and citation resolvability. The exercise cost \$8.77 and fetched 120 pages.

Two automated country assessments. The pipeline researches all fifty-seven indicator rows and the two carried candidates for one country, proposes a source and a tier for each, sets a level, and records a rung-by-rung argument including the negative finding for every level below five. It then runs a second review on the rows where being wrong costs the most, using a different vendor, and applies whatever survives. Egypt cost \$21.88 across both passes and took 38 minutes. Nigeria cost \$21.47 and took 37 minutes.

A comparison against the verified assessments. Every automated row was compared with its counterpart in your assessments. The pipeline did not read those assessments at any point, and the comparison does not write to them.

It is important to point out what these figures do not establish. Two countries is a small base, one vendor did the first pass in both, and each country was assessed once. Two independent runs of the Egypt first pass agree on 53 of 57 rows, so a difference of one or two rows anywhere below is inside the noise and only larger patterns should be read as findings.

Results in Summary

	Egypt	Nigeria
Rows at the level your assessment records (of 57)	24	30
Prerequisites at the same status (of 12)	5	5
Rows read above your assessment	2	4
Levels withheld where your assessment set one	21	19

	Egypt	Nigeria
Recorded gaps (your assessment)	20 (5)	22 (4)
Cost, both passes	\$21.88	\$21.47

2. Indicator Definitions (Decision 13.5)

Section 13 calls this the largest item, with 44 of 57 rows carrying an open definitional question and 8 of the 12 prerequisites among them. The automated runs give that item a measurable consequence, and it is the finding in this package most worth your time.

The case for leaving the definitions as they are is real and should be stated first. Each of the 44 questions was raised by the instrument's own audit rather than by a reader, the proxies chosen are defensible, and the two hand assessments were completed without any of them being resolved. An assessor who knows the sector applies a sensible reading and moves on, and thirty of the forty-four are recorded as a defensible but measurably different proxy rather than as a construct error.

A machine does not have that judgment, and this is where the cost becomes visible. **Every one of the six rows where the automated pipeline read above your assessment, in both countries, carries an open definitional question, and the instrument's own audit classifies all six as construct drift.** Not one of them is a case of the machine misreading a source. In each, the machine took the name as written and read it permissively, and the assessors read the same name restrictively.

Device financing and subsidy schemes (Egypt, 2.11). Your assessment records level 1 and states that no government or telecommunications operator scheme targeting low-income or rural users was found. The machine recorded level 3, citing the ProGIG initiative run by Egypt's National Telecommunication Institute with Nasser Social Bank, which offers soft loans for the purchase of laptops. The name fixes neither a targeting test nor whether a subsidy is required. The assessor applied a targeting test and found nothing; the machine applied none and found something.

Responsible-AI safeguards, consent and rights (Egypt, 7.12). Your assessment records level 1 and states that no agriculture-specific consent or farmer-data-rights safeguards exist beyond the general regime. The machine recorded level 3, citing the Egyptian Charter for Responsible AI issued by the National Council for Artificial Intelligence in 2023. The assessor required the safeguards the name specifies; the machine accepted a general instrument that does not use the words consent or rights.

Inter-ministerial coordination mechanism (Nigeria, 4.9). Your assessment records level 1. The machine recorded level 3 on a United Nations record of a Nigerian inter-ministerial committee, reading the name as satisfied by a committee that exists and meets. The assessor read it as requiring a mechanism that coordinates.

The second review does not correct this class of error, and it is worth being clear about why. On Nigeria's farmer registry (3.3) it did: the first pass recorded level 3 and the review returned it to

your level 2, after reading the Ministry's own description of a November 2025 event as a stakeholder workshop written in prospective language. But 4.9 above was created by that same review, which filled a gap the first pass had left. The reviewer reads sources more carefully than the first pass. It reads the same unratiſed names, and gains nothing on them.

The pattern is consistent across both countries and points in one direction. A permissive reading of an unratiſed name produces a higher level, higher levels move prerequisites, and prerequisites move whole columns of the readiness matrix. In this sense the definitions are not a documentation item. They are the mechanism by which an automated assessment becomes either conservative or generous, and at present the model does not specify which.

3. The Binding Rules and the Tier Bar (Decisions 13.4 and 13.8)

One row moved every column of Nigeria's readiness matrix, and the ruling that would prevent it sits across two of the twelve items.

Your Nigeria assessment reads Partial on five use-case columns and Blocked on the sixth. The automated assessment reads **Unverified on all six**. The cause is a single row: rural mobile broadband coverage (2.1), which your assessment records as Present at level 3 and the machine withheld. Binding rule *universal-unverified* then returns every column Unverified, on the stated ground that a known blocker outranks an unknown one. The rule did exactly what it says. The question is whether what it says is what you want when the levels are machine-set.

The reason the machine withheld 2.1 is the more consequential half. Your assessment records rural 4G population coverage at 48 percent for 2024, sourced to GSMA geospatial analysis, at tier 4. The pipeline requires tier 1 to tier 3 quote-verified evidence for a prerequisite, on the ground that a one-level error in a prerequisite propagates to six columns. **The automated pipeline therefore could not have reproduced your Nigeria 2.1 level at any depth of searching.** The bar would have withheld it whatever the machine found, because the only published estimate is modelled and modelled estimates sit at tier 4.

Two readings are available and the choice is yours. Either the bar is right and a universal prerequisite resting on a modelled tier 4 estimate should read Unverified, in which case Nigeria's matrix reading is a correction rather than a loss. Or the bar is too high for an indicator that no one measures directly, in which case the tier assignment in decision 13.8 needs to say so explicitly, because a general rule that tier 4 cannot carry a prerequisite will keep producing this result in every country where rural coverage is modelled rather than measured. That is most of them.

Egypt shows the same rule operating with no visible effect, because its 2.1 is a recorded gap in both assessments and its matrix already read Unverified. The rule is not new. What is new is that an automated pass reaches the withheld state far more often than a hand pass does, which turns a rule that rarely fired into one that fires by default.

4. Quality Control and the Scope of Gate 2 (Decision 13.9)

Section 14 records an amendment to the QC Protocol: peer review must ask whether the evidence answers *this* indicator, not only whether the source states the number. The audition provides direct evidence for that amendment, and the evidence is stronger than expected.

Quote verification, which checks that every recorded value appears verbatim in the page cited, caught invention completely. Across thirteen cells and five models there was one fabricated citation, from a single model, and four of the five models produced none.

Quote verification did not catch construct substitution at all. On the Nigeria digital ID row, one model recorded 121,191,781, a raw count of National Identification Number enrolments, against an indicator whose name ends in “(%)”. On the Nigeria mobile money row, the same model recorded 47.1, an unweighted frequency drawn from a thousand-case microdata file, from a page that carries an explicit warning that such counts cannot be interpreted as summary statistics. Both figures were verbatim on the pages cited, and both passed quote verification without difficulty. Four of the five models named the warning and declined to record a value.

In short, the check that stops fabrication and the check that stops the wrong number being recorded are different checks, and only the second is a construct test. The section 14 amendment is what closes that gap, and the audition suggests it should be ratified as written rather than softened.

One figure bears on the scope of Gate 2. The design work assumed a second review would cover roughly twenty rows. Scoped as intended, to prerequisites, withheld levels, and recorded gaps, it covered 38 rows in Egypt and 36 in Nigeria, because an automated first pass withholds more often than a hand one. It cost \$6.39 and \$6.14 respectively, against an allocation of \$75, and changed three levels in each country. It is cheap and it earns its place, but its scope is a function of how conservative the first pass is, and that in turn is a function of section 13.5.

5. The Source-Tier Lookup (Decision 13.8)

The starter domain-to-tier lookup in the Source-Tier Protocol has been implemented as code, which means it can now be applied identically to every country and every row rather than remembered. Implementing it surfaced two assignments the written protocol does not resolve, both of which produced wrong answers before they were fixed by hand.

The first is the World Bank’s own domains. Analytical reports and flagships are published at `openknowledge.worldbank.org`, which is a subdomain of the same host that serves the statistical databases. A lookup keyed on the shorter name files a flagship report as an official statistic. The second is the United Nations. `publicadministration.un.org` serves the E-Government Knowledgebase, a statistical database, while `news.un.org` serves the organization’s newswire, and the protocol’s tiers put those four tiers apart.

Both were resolved provisionally in code, the first as tier 2 and the second as tier 1 with the newswire held at tier 5. Ratifying them, together with the two open judgment calls the protocol already flags on national statistical offices with documented quality concerns and on donor

project completion reports, would make the lookup deterministic. The practical value of that is not tidiness. A ratified lookup is executable, and an executable lookup applies the same tier to the same publisher in Cairo and in Abuja without an assessor having to remember what was decided last time.

6. Where the Runs Are Silent

Six of the twelve decisions gained no evidence from this work and should be read as untouched by it: band edges (13.1), A1 thresholds (13.6), sub-readings display (13.7), the Practice Library schema (13.10), Bhutan's exclusion (13.11), and the per-use-case prerequisite mapping (13.3), which the runs applied without testing.

Two gained something slight. On the A1 candidates (13.2), the machine located Egypt's irrigation figure at 100 percent of cultivated area against the 94.21 percent your assessment records from AQUASTAT for 2023, and found neither candidate for Nigeria; the candidates are findable but not reliably so. On whether need and outcome indicators belong in a readiness mean (13.12), the automated runs withhold enough levels that use-case means rest on fewer rows than in your assessments, which makes any status driven by the mean rather than by a prerequisite thinner than it looks. Neither observation is strong enough to argue from.

7. What This Does Not Claim

The automated assessments are test output. They have not been reviewed, they are not proposed for any use, and they should not be read as second opinions on Egypt or Nigeria. Where they disagree with your assessments the presumption is that your assessments are right, and the value of the disagreement is diagnostic.

The verified assessments are the best available answer rather than the truth, and the comparison is calibrated against them for want of anything better. The pipeline has not been run on any country outside the two that already have hand assessments, and it will not be until the section 13 rulings land.

Lastly, it is important to mention that the abstention behaviour described throughout this note is a design choice and not an accident. The pipeline is built to withhold a level rather than to guess one, on the ground that a withheld level is visible and a guessed one is not. Roughly two thirds of the rows that both countries withheld were rows whose evidence the pipeline never reached, rather than rows it read and declined to score. Those are a retrieval problem with named sources and a known fix. They are not an argument for relaxing the rule.

8. The Decision That Releases the Rest

The engineering is finished and measured. What it is waiting on is not more engineering.

Every automated assessment produced so far rests on 44 indicator definitions that have not been ratified, and every material disagreement between the machine and your assessors comes from that fact rather than from the evidence, the vendors, or the cost. A ruling on section 13.5, taken with the tier assignment in 13.8 and the binding rules in 13.4, determines whether an automated national assessment reads conservatively or generously, and it determines it once for every country the instrument is ever applied to. Extensive commitment at the point of ruling is what converts a working pipeline into one that can be trusted with a country that has no hand assessment to check it against.