Vendor decision — primary and independent second

24 August 2026. Taken from the audition in `AUDITION-RESULTS.md`, which is the measurement; this file is the judgment made on it. The audition can be re-scored and re-run without disturbing this record, and this record should be re-taken if it is.

The decision

Role	Vendor	Why
Primary research vendor	`anthropic/claude-opus-5`	Tied at the top on all three governing scores; widest discovery, surfacing 64 pages of which 46 were T1–T3, and an admissible page in 12 of 13 cells
Independent second, for Gate 2	`openai/gpt-5.6-terra`	A genuinely different vendor, tied with the primary on all three scores, and the strongest of the three gpt-5.6 siblings on level accuracy and tier compliance
Discovery peer	`perplexity/sonar-pro`	Confirmed in the role decision C6 assigns it and no other — see below
Not selected	`gemini/gemini-3.1-pro-preview`	The only entrant that fabricated, and the only one that recorded two figures measuring something other than what the indicator names

The primary is a near-tie, and should be recorded as one. `anthropic/claude-opus-5`, `openai/gpt-5.6-terra` and `openai/gpt-5.6-sol` returned identical figures on fabrication (0%), tier compliance (100%) and citation resolvability (100%), and identical value and level accuracy (60%). Nothing in the three governing scores separates them. Discovery breadth is the tie-breaker, and it is a weaker measurement than the three — so if the primary needs to be swapped later, this decision should not be treated as an obstacle.

Cost was deliberately not used as a tie-breaker. The Anthropic rates in `prices.json` are published; the OpenAI and Gemini rates are placeholders set at Opus-tier so the counter cannot read low. Comparing a real price against a placeholder would be a measurement of the placeholder. The recorded token counts make the comparison exact the moment the real rates are entered.

The one tier failure left standing is real, and it errs safely. After the two scoring corrections, a single non-compliance survives: `gpt-5.6-luna` cited the UN E-Government Knowledgebase — the correct source, correctly quoted, correctly valued — and proposed **T5** for it. That is an official UN statistical database being filed as news and vendor material. Under-tiering never inflates a claim; what it does under this pipeline's own rules is worse than harmless and better than the alternative — a T5 citation cannot yield `Documented`, so the row would fall to `Judged`, and on a prerequisite it would fail the C4 bar and be held. The failure mode is lost coverage, not false confidence. It is still the reason terra and sol are preferred within the family.

Independence is by vendor, not by model. The three gpt-5.6 siblings behaved almost identically across all thirteen cells — they abstained on the same four known cells and detected all three absences — which is what one would expect of models sharing a lineage. A Gate 2 run on a sibling of the primary would refute very little, so the second is drawn from a different vendor even though a sibling scored the same.

What the audition established

A fabrication baseline now exists, and it is zero. Four of the five entrants asserted no value across thirteen cells that they could not quote verbatim from a page they had been given. This is the number a reviewer of a machine-drafted roadmap will ask for and which could not be stated before today. It is a baseline *under quote verification against a known page set*, which is the condition the pipeline itself enforces — not a claim about unconstrained generation.

All five entrants detected all three non-existent constructs. Including **N1** — recording Egypt's national 99.8% mobile broadband coverage against an indicator naming *rural* coverage. That error passed a hand assessor gate and an initial peer review in the verified Egypt assessment before an audit caught it; it is defect #1 in the issues log. No entrant repeated it. Every one of them reported the national figure as context and refused it as an answer.

The binding constraint is retrieval depth, not judgment. On four of the ten known cells (K6, K8, K9, K10) every entrant abstained, and in each case the reason was that the page carrying the answer was not among the ten they were given:

- **K6** — the ITU price basket lives behind a JavaScript dashboard that returned navigation text and no values.
- **K8** — the supplied pages carried absolute NIN enrolment counts with no population denominator, so a percentage could not be quote-verified.
- **K9** — the rural smartphone figure sits in a national household survey; the supplied pages offered national ownership and rural mobile-internet *use*, neither of which is the named construct.
- **K10** — no supplied page carried a published population estimate.

So the measured picture is a fabrication rate near zero and an abstention rate on known-answer cells of 40%. That asymmetry is the one to act on: the system's error mode is silence, not invention, and more retrieval depth buys more than more model.

Quote verification catches invention. It does not catch construct substitution. This is the sharpest finding for the pipeline's design. On K10 the supplied Findex page carried the number 47.1 — an unweighted frequency from a 1,000-case microdata file, beside the page's own warning that such counts "cannot be interpreted as summary statistics". Four entrants named that warning and abstained. One recorded 47.1 as Nigeria's mobile money account rate, fully quote-verified. On K8 the same entrant recorded 121,191,781 — a raw enrolment count — against an indicator whose name ends in "(%)", again fully quote-verified.

Neither would have been caught by the quote gate, because neither was invented. Both are exactly what decisions **C3** (abstain where the evidence measures a different construct) and **C4** (prerequisites need T1–T3 quote-verified evidence) exist to stop, and they are the reason those gates cannot be simplified away into quote verification alone.

Perplexity earned its assigned role and no more. Its citations were the most efficient of any proposer — 25 of 29 pages it surfaced were T1–T3, the highest ratio in the audition — but it reached an admissible page in only 10 of 13 cells against 12 for the leading entrants. That is a good peer and a poor primary, which is precisely the position decision C6 assigns it: its citations are re-fetched through Jina and quote-verified, and its prose never becomes a source of record.

What this decision does not settle

- **Whether Gate 2 earns its 15%.** That needs the second vendor actually running a refutation pass, which is Thread 3. The shadow-run figures recorded now are the before-measurement for that comparison.
- **The abstention threshold.** The audition measures abstention against thirteen cells where the answer is known. The shadow run measures it against fifty-seven where the whole assessment is known, and that is the calibration that matters.
- **Real cost ranking.** Pending the OpenAI and Gemini rates in `prices.json`.