

Gate 2 — second review of EGY_shadow (Egypt)

An independent vendor (openai/gpt-5.6-terra) attacked the first pass on the **38 rows** where being wrong costs the most: the twelve prerequisites, every ratification hold and every recorded gap. It planned its own searches from the construct and never saw the first pass’s sources while searching; it saw the recorded row only when forming its verdict. Its proposals entered through the same gates as any other evidence.

\$6.39 across 646 vendor calls in 14 minutes.

What it changed

outcome	rows	
filled	3	a gap or withheld level replaced by evidence that cleared every gate
withdrawn	0	a level removed: the cited source was shown not to support it
relevelled	0	the level changed: a second reading of better evidence, cleared through every gate
adjusted	2	substance and level stand, provenance corrected
upheld	33	the row survived the attack

Verdicts as the reviewer gave them

verdict	rows
confirmed	24
adjust	9
refuted	5

A reviewer verdict is not the same thing as a change to the row. A refutation that proposes evidence still has to clear quote verification, the tier rule, the construct rule and — on a prerequisite — the T1–T3 bar, exactly as the first pass did. Where the two columns disagree, the gates are why.

Rows Gate 2 filled

Evidence the first pass did not reach, which then cleared every gate. This is the measure of what a second independent search buys.

3.7 Satellite/EO data integration · hold · severity medium

- The first pass's withheld conclusion is too broad. The ESA record is Egypt-specific evidence of a completed EO service demonstration designed to implement EO-based operational monitoring capacity and identifies an Egyptian institutional user. It does not justify an Operating rung, but it does support Adopted under the indicator name as written.
- now: Documented at level 3, T3 — ESA Data User Element — TIGER Innovators - Satellite Monitoring of Lake Water Quality in Egypt
- https://due.esrin.esa.int/page_project79.php

4.9 Inter-ministerial coordination mechanism · prerequisite · severity high

- The first-pass Gap is contradicted by supplied T1 evidence: Source 2 explicitly records an Egyptian Inter-Ministerial Committee formed by a prime-ministerial decree. This is affirmative evidence of a formally adopted national coordination mechanism, although it does not establish the committee's ongoing operation or performance.
- now: Documented at level 3, T1 — United Nations Statistics Division, National Monitoring of the SDGs
- https://unstats.un.org/sdgs/files/meetings/sdg-inter-workshop-june-2018/Day1_Session2_Egypt.pdf

8.4 Mobile money account (%) · gap · severity high

- The first pass's gap is refuted by the supplied World Bank Global Findex DataBank table, which publishes the exact Egypt mobile-money-account series and shows 0.0 for 2021.
- now: Measured at level 1, T1 — World Bank, Global Findex database - DataBank
- <https://databank.worldbank.org/id/739ae3c1>

Provenance Gate 2 corrected

The substance and the level survived; the source, tier or vintage did not.

4.1 Data protection/privacy law · prerequisite · severity low

- The T3 source verifies promulgation and identifies Law No. 151 of 2020, supporting Adopted and level 3. However, the first-pass value states that the law was effective on 15 October 2020, whereas the cited page—dated 10/14/2020—says only, “The law will become effective on October 15, 2020.” The source therefore does not quote-verify the first pass's past-tense effective-status assertion, although this correction does not change the rung or level.

5.5 Digital extension capability · prerequisite · severity medium

- The first pass correctly identified an Egypt-specific digital extension artefact and correctly left quality and scale blank, but its Adopted rung is too low. The independently supplied FAO country case affirmatively states that a campaign was launched to expand use of the Egyptian app, establishing an operating service rather than only a

launched/adopted design. The mechanical level remains 3 because Operating without quality or scale evidence scores 3.

Every row reviewed

id	role	reviewer verdict	outcome	what it found
1.5	hold	confirmed	upheld	The first-pass withholding survives the attack. The additional supplied Ministry material does not supply an a
1.6	gap	confirmed	upheld	The gap survives the attack. The strongest potentially relevant pages are the Ministry PRIME intervention reco
1.7	gap	confirmed	upheld	The first-pass gap survives the supplied-source attack. A potentially relevant Global Findex construct is expl
1.8	gap	confirmed	upheld	The first-pass gap survives review. The strongest potentially relevant supplied source is the 2026 peer-review
2.1	prerequisite	confirmed	upheld	The first-pass gap survives review. Source 1 affirmatively supports the claimed

id	role	reviewer verdict	outcome	what it found
2.7	gap	adjust · tier or class wrong	upheld	construct mismatch: it labels The first pass's substantive withholding of a level holds, but its gap search missed a published rural househo
2.9	prerequisite	confirmed	upheld	The first-pass value, vintage, T1 publisher tier, and rural construct all survive review: Source 1 directly di
3.3	prerequisite	confirmed	upheld	The first-pass withholding survives the adversarial check for the indicator as written. A previously unreached
3.4	hold	confirmed	upheld	The withholding survives the independent attack. No supplied T1–T4 page states a fact showing digital land/plo
3.5	hold	confirmed	upheld	The withholding was attacked with the retrieved Egypt-specific Parquet/CC-BY dataset. That is an existence fac

id	role	reviewer verdict	outcome	what it found
3.7	hold	adjust · better evidence exists	filled	The first pass's withheld conclusion is too broad. The ESA record is Egypt-specific evidence of a completed EO
3.8	hold	refuted · construct mismatch	upheld	The first pass incorrectly imposed the unresolved, narrower interpretation of an operating digital soil-inform
3.11	prerequisite	confirmed	upheld	The first-pass gap survives review: no supplied admissible source quote-verifiably establishes an agricultural
4.1	prerequisite	adjust · source does not support the value	adjusted	The T3 source verifies promulgation and identifies Law No. 151 of 2020, supporting Adopted and level 3. However
4.4	hold	confirmed	upheld	The first-pass construct hold survives the supplied-source attack. No page supplies affirmative evidence of a
4.5	prerequisite	confirmed	upheld	The first-pass gap withstands

id	role	reviewer verdict	outcome	what it found
4.7	prerequisite	adjust · better evidence exists	upheld	review. The potentially relevant supplied evidence does not meet the named const The first pass correctly withheld a percentage-based level, but its evidence class should not remain a pure ga
4.9	prerequisite	refuted · better evidence exists	filled	The first-pass Gap is contradicted by supplied T1 evidence: Source 2 explicitly records an Egyptian Inter-Mini
5.3	gap	confirmed	upheld	The first pass's claim that no source supports the named STEM construct survives review. A T1 Egypt value is n
5.4	gap	confirmed	upheld	The independently surfaced T5 deck is the closest additional lead, but its reported ICT/IoT knowledge measure
5.5	prerequisite	adjust · tier or class wrong	adjusted	The first pass correctly identified an Egypt-specific

id	role	reviewer verdict	outcome	what it found
5.7	prerequisite	confirmed · could not locate independently	upheld	digital extension artefact and correctly left quality an The supplied sources do not affirmatively refute the first pass's inaccessible CGIAR/IFPRI citation, so failur
5.8	gap	confirmed	upheld	The first-pass gap survives the supplied evidence. The potentially refuting MCIT GreenHack source is affirmati
5.12	gap	confirmed	upheld	The first-pass gap survives review of all supplied pages. None states a value for gender balance in Egypt's di
6.3	gap	confirmed	upheld	The first-pass gap survives review. The official World Bank Egypt- specific B- READY topic pages visibly contain
6.4	hold	adjust · tier or class wrong	upheld	The first pass's stated basis for withholding —"the only source is T5"— does not hold:

id	role	reviewer verdict	outcome	what it found
6.9	gap	confirmed · could not locate independently	upheld	Source 1 is a T2 CGIAR co The first-pass gap survives the supplied-corpus attack. The closest affirmative lead is Source 10, but it is a
6.12	hold	adjust · tier or class wrong	upheld	The first pass's withholding rationale misapplies the T5 rule. T5 material may establish existence facts, and
6.13	gap	adjust · better evidence exists	upheld	The first pass correctly found no usable percentage, but its Gap classification overlooks the supplied ICARDA
6.14	prerequisite	confirmed	upheld	The first-pass withholding survives review. The supplied FAO T1 page affirmatively supports the narrower point
7.2	hold	refuted · better evidence exists	upheld	The first-pass withholding rationale fails because it did not reach the official MCIT release. That release ex

id	role	reviewer verdict	outcome	what it found
7.12	prerequisite	refuted · construct mismatch	upheld	The first-pass T3 citation does support the existence and launch of a general responsible-AI charter, but it d
8.4	gap	refuted · better evidence exists	filled	The first pass's gap is refuted by the supplied World Bank Global Findex DataBank table, which publishes the e
8.5	hold	confirmed · could not locate independently	upheld	I attacked the withholding rationale by looking for an Egyptian sex-disaggregated land-ownership statistic in
8.9	gap	confirmed	upheld	The gap was attacked using the supplied Egypt-specific ICARDA and IFPRI materials, especially the smallholder
8.11	gap	adjust · better evidence exists	upheld	The first-pass conclusion that no Egyptian local-language service evidence exists is too broad: the Egyptian T
8.12	gap	confirmed	upheld	The gap survives the supplied-page attack. The closest Egypt-

id	role	reviewer verdict	outcome	what it found
8.17	gap	confirmed	upheld	specific impact study measures raised-bed technol The supplied pages do not affirmatively refute the recorded gap: none quote- verifiably reports a national or o

Reading this

Gate 2 is scoped to the rows the first pass was least sure of, so a high upheld count is not a null result — it is the first pass’s abstentions being confirmed as real by an independent searcher, which is exactly what makes a recorded gap trustworthy. The figure that decides whether Gate 2 earns its share of the budget is **filled** plus **withdrawn**: rows where a second opinion changed what the assessment says.